

MVProc is a Model-View-Controller web framework that uses MySQL stored procedures as the controller element. It is implemented as an Apache2 module, mod_mvproc.

There are currently no plans to implement this for any other database, as no other database has enough flexibility of functionality in the context of stored routines. (At least as far as I'm aware.)

Installation

The Makefile and #includes section of mvproc.h may need to be edited before compiling, specifically the APACHE_HOME variable and location of apr-1 and apreq2. It will build on Ubuntu, but other distros may have different locations of headers and libs. If you have difficulties regarding libuuid.la missing, you'll have to build libapreq2 from source.

```
(apt-get source libapreq2;)
```

and configure it to use the available apache installation:

```
(configure --prefix=/usr --with-apache2-apxs=/usr/bin/apxs2;)
```

Configuration

Since `mod_mvproc` is a module it must be loaded in your apache configuration file. `apreq_module` is required and must be loaded prior to loading `mvproc`, like so:

```
LoadModule apreq_module /usr/lib/apache2/modules/mod_apreq2.so
LoadModule mvproc_module /usr/lib/apache2/modules/mod_mvproc.so
```

There are eleven (11) configuration directives for `mod_mvproc`:

mvprocSession (Y or N) – sets session behavior. If sessions are enabled (Y), the module will supply called stored procedures with a user variable `@mvp_session`, which will contain a 32 character session id. A cookie will be produced and maintained with this id as its value. The procedure may set the `@mvp_session` value, which will update the value of the cookie. The value is also included in the output of the call in the `PROC_OUT` section.

mvprocDbGroup (yourDBgroup) – this is the group in mysql's `my.cnf` file which `mvproc` will use to authenticate and connect to the database. An example:

```
[mydb]
host      = localhost
port      = 3306
user      = myuser
password  = mypass
database  = webdb
```

Please refer to the MySQL documentation for `my.cnf` subtleties.

mvprocTemplateDir (/your/template/directory) – the location of the templates. This is best placed outside the doc root – in fact, if the template files are in the doc root Apache will get confused, thinking it should serve the file instead of allowing `mvproc` to handle the transaction. If `mvprocTemplateDir` is not set, or set with an empty string, templates and layouts will not be used. More on templates later.

mvprocDefaultLayout (yourLayoutTemplate – leave off the '.tpl') – If set, this will set the `@mvp_layout` session variable to a default value. It's not required to set this in the configuration, as procedures can set it as well. More on layouts later.

mvprocCache (Y or N) – Y enables caching of procs AND templates. Set to N during development and Y for production use. Also available are T – cache templates only, and D – cache database only.

mvprocDbPoolSize (number) – The number of connections to pool. If this is not set, or set to zero, a new connection will be established, used, and closed for each request.

mvprocOutputStyle (MIXED | PLAIN | JSON | EASY_XML | EASY_JSON) – see the Output section below.

mvprocErrTemplate – If `mvprocTemplateDir` is configured, this is the name of the template to render when there's a database error. The error can be accessed as `<# status.error #>`.

mvprocAllowSetContent (Y or N) – provides `@mvp_content_type` session variable, which allows a procedure to tell Apache how to set the Content-Type header.

mvprocAllowHTMLfromDB (Y or N) – default N, Y allows HTML output from the database to be rendered as HTML. By default, MVProc will convert '<' to '<'; '&' to '"'; etc. to prevent XSS and CSRF attacks.

mvprocUploadDirectory - Sets the path to a directory where uploaded files are written. It is highly recommended that this path be OUTSIDE the docroot. If the path doesn't exist, MVProc will attempt to create it. Default value is /tmp

Requests

Requests are typical “pretty” web format, like so: <http://mysite.com/procname>. If the url points to a file that exists, the module will decline handling, which means the file will be served as normal. If the url is requested without a procedure specified, a procedure named 'landing' will be looked for. If a procedure does not exist with the requested name, the module will decline processing.

GET and POST requests are supported, as well as file uploads. A file upload will be written to the system's temp directory with a pseudo-random name plus the original file extension. This filename will be reported to the procedure through the uploaded file's variable name. So

```
<input type="file" name="myfile">
```

will send something like '/tmp/F228FA.jpg' into the procedure as the IN or INOUT argument 'myfile', but ONLY IF THE ARGUMENT EXISTS in the procedure's definition.

Procedures

All aspects of MySQL stored procedures are supported. IN, INOUT, OUT parameters and multiple result sets are available to the template parser and are shown in xml or json output. User variables available to procedures are currently:

- `@mvp_servername` – the server hostname (mysite.com)
- `@mvp_requestmethod` – GET, POST, or REQUEST
- `@mvp_uri` – the unparsed uri of the request
- `@mvp_template` – by default, the procedure name – this value can be changed in the procedure to tell the module to use a different template file. This can be 'procname' or 'myother_template' or even (under the template directory) 'my_subdirectory/my_other_sub/even_further/as_deep_as_you_like/template_name'
Remember to leave off the '.tpl'
- `@mvp_layout` – the layout template to use – this value can be set or changed in the procedure. Handling is essentially the same as for `@mvp_template`.
- `@mvp_headers` – the HTTP headers
- `@mvp_remoteip` – the requestor's ip address
- `@mvp_session` – (if `mvprocSession = 'Y'`) the current value of the MVPSESSION cookie
This can be set in the procedure, but the module will set it up and provide a value by default. Once overridden, the set value will be returned with each subsequent request. (Just like a session... heh.) You could even store data in the cookie if it's ok for everyone to see. The only size limit is what a browser will actually send and receive in a cookie.
- `@mvp_content_type` – (if `mvprocAllowSetContent = 'Y'`) – This is passed in as an empty string. It's only used in output if set by the procedure.

Known Issues

Returning multiple result sets with the same table name causes libmysqlclient to seg fault. This includes calculated selects. So if you want to do this:

```
SELECT 'a value' AS one_value; SELECT 'more value' AS too_value;
```

Instead, do this:

```
SELECT 'a value' AS one_value, 'more value' AS too_value;
```

And if you want to return multiple result sets from the same table, use aliases.

If libmysqlclient starts to allow multiple result sets from the same table, be cautious about using EASY_XML and EASY_JSON output. The results are, at this point, undefined.

Output

XML MIXED is the default output of a request. If no template exists, the module goes with this. Result sets are rendered in xml as a <table> element with child <row> elements.

For the MIXED output type, columns whose declared size is 32 or less (for example VARCHAR(24)) are shown as attributes of rows, while blobs and columns of greater size (like VARBINARY(4096) or CHAR(65)) are shown as children of rows with CDATA encapsulated values.

The PLAIN output type populates the name attribute of the table tags, but all columns are rendered as CDATA encapsulated child elements.

Another option is JSON, which outputs an array of objects named 'table', each of which owns a name value and an array of row objects each of which owns one named value per column. Most ajax libraries support this format.

EASY_XML is the PLAIN type, except that instead of <table name="tableName"> elements, the output is written like <tableName>. This may cause problems if multiple result sets from the same table are output.

EASY_JSON is JSON, except that instead of an array of objects named 'table', the output is an object with each result set as an element named for its table. This may cause problems if multiple result sets from the same table are output.

All result sets are output, as well as a "table" called PROC_OUT, which holds all INOUT and OUT values as well as mvp_session and mvp_template values. If a table is not declared in a select out (eg. SELECT 'whatever' AS myval) the table will be called "status".

Here's an example of MIXED:

```
<results>
  <table name="widgets">
    <row idwidget="1" label="demo" version="1.00">
      <description>A demo for testing</description>
    </row>
  </table>
  <table name="PROC_OUT">
    <row>
      <mvp_session>19e820417390d3bf564261886d70ea64</mvp_session>
      <mvp_template>getWidgets</mvp_template>
    </row>
  </table>
</results>
```

Errors will look like this:

```
<results>
  <table name="status">
    <row>
      <error>Please consult your documentation for correct usage.</error>
    </row>
  </table>
</results>
```

Templates

Templates are typically html pages with tags for the parser to use to plug in values, loop through result sets, conditionally show or not show markup, include other templates. Some simple rules:

- an MVProc tag looks like this: `<# the tag is in here #>`.
- String literals are between single quotes.
- Inside single quotes you can have any characters you want – the quotes escape everything.
- Escape single quotes within single quotes with a backslash: `'\'`
- The tag names must be either ALL CAPS (ELSEIF) or all lower (elseif).

The template tags are intentionally few in number. It encourages separation of concern.

- **Value** - `<# [table.]field_name[[row_num]] #>`
The default table is PROC_OUT, and the default row_num is 0. Inside a LOOP, the default table becomes the LOOP table and the default row becomes the LOOP's CURRENT_ROW.
CURRENT_ROW (all caps) is a built-in value of type INT (actually unsigned long in C terms).
NUM_ROWS (all caps) is a built-in value of type INT accessed like: `<# IF tableName.NUM_ROWS > 0 #>`
The '@' table holds user variables for the templates (see SET).
- **IF** - `<# IF myvar = 'hello' #>`
Supported comparison operators are: `=`, `==`, `!=`, `!`, `<>`, `<`, `>`, `<=`, and `>=`
With no operator, a value equals true if non-zero (int & float) or non-empty (string).
The not operator (!) equals true if zero (int & float) or empty (string).
Nesting is supported, as well as AND, and, OR, or, &&, and ||
AND, and, and && take precedence over OR, or, and ||
Example: `<# IF mytable.what[2] = 'ok' AND (CURRENT_ROW > 4 OR !@.checkit) #>`
Nesting is arbitrarily limited to a depth of 64.
No math is supported in IF tags.
Constants are supported: strings are quoted with single quotes and floats must have a '.' (eg 0.0)
- **ELSIF** - `<# ELSIF @.val_is_set #>` - Identical parsing and evaluation to IF.
- **ELSE** - `<# ELSE #>` - This functions as anyone would expect.
- **ENDIF** - `<# ENDIF #>` - Again, like anyone would expect. This tag is REQUIRED with IF usage.
- **LOOP** - `<# LOOP mytable #>`
The LOOP tag begins a template segment that will iterate once for each row in a result set.
Inside the LOOP, the table specified becomes the default table, and CURRENT_ROW will evaluate to the current row (zero indexed).
- **ENDLOOP** - `<# ENDLOOP #>` - Closes a loop. This tag is REQUIRED for each LOOP started.
- **INCLUDE** - `<# INCLUDE another_template #>`
The specified template will be included at the tag's position.
The included template is referenced from the configured myprocTemplateDir, so if your template directory is /var/templates and you `<# INCLUDE layouts/header #>`, the file looked for will be /var/templates/layouts/header.tpl
Always leave the '.tpl' off the include argument.
INCLUDE can also accept a value tag like: `<# INCLUDE myTable.myValue[4] #>`
- **TEMPLATE** - `<# TEMPLATE #>` - This is essentially an include tag, except that it uses the @mvp_template session variable. Use this convenience tag inside a layout template.
- **SET** - `<# SET myvar = 'ok' #>`
One could create a very rich site with lots of functionality without using this tag. I haven't benchmarked the difference between setting all required values in the procedure vs. setting some in the template but I suspect SET might be a bit slower, and it's certainly less efficient with memory. That being said, is it available.
The SET tag sets a value in the '@' table, which supports only one row.
`<# SET row_class = 'color' + CURRENT_ROW % 2 + 1 #>` would be referenced
`<# @.row_class #>` (and this is useful for alternating row css style)
Supported operators are `=`, `+`, `-`, `*`, `/`, `%`, and `comma(,)`
String "math" supports only concatenation (+), ints and floats use C-style math.
`*`, `/`, and `%` take precedence over `+` and `-`.
Modulus (%) evaluates to the remainder for int values and the fraction for floats.
Constants are supported: strings are quoted with single quotes and floats must have a '.' (eg 0.0)
The comma is for multiple assignments in a tag, like so: `<# SET a = 0, b = 'hi', c = @.a / 1 #>`

And that's it.

I think. Any questions, bugs, requests – please post on sourceforge.net

-Jeff Walter
maintainer, MVProc